

# Operating System

**Unit – 3**

**(Part-B)**

## Process Synchronization




**Mayank Mishra**

**School of Electronics Engineering**

**KIIT-Deemed to be University**

# Mutex Locks

- The hardware-based solutions to the critical-section problem are complicated as well as generally inaccessible to application programmers.
- Instead, operating-system designers build higher-level software tools to solve the critical-section problem. The simplest of these tools is the mutex lock. (In fact, the term **mutex** is short for **mutual exclusion**.)
- Mutex Lock is used to protect critical sections and thus prevent race conditions. That is, a process must acquire the lock before entering a critical section; it releases the lock when it exits the critical section.
- The *acquire()* function acquires the lock, and the *release()* function releases the lock, as illustrated in Figure.



*Fig. Solution to the critical-section problem using mutex locks*

## Mutex Locks (Contd.)

- A mutex lock has a boolean variable `available` whose value indicates if the lock is available or not.
- If the lock is available, a call to **`acquire()`** succeeds, and the lock is then considered unavailable. A process that attempts to acquire an unavailable lock is blocked until the lock is released.

The definition of `acquire()` is as follows:

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}
```

The definition of `release()` is as follows:

```
release() {  
    available = true;  
}
```

# Mutex Locks (Contd.)

## Process P1

The definition of acquire() is as follows:

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}
```

The definition of release() is as follows:

```
release() {  
    available = true;  
}
```

## Process P2

The definition of acquire() is as follows:

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}
```

The definition of release() is as follows:


```
release() {  
    available = true;  
}
```

- Calls to either **acquire()** or **release()** must be performed atomically.
- *An atomic operation in an operating system (OS) is a sequence of instructions that are executed as a single unit without interruption.*
- The main disadvantage of the implementation given here is that it requires **busy waiting**.
- While a process is in its critical section, any other process that tries to enter its critical section must loop continuously in the call to **acquire()**. This continual looping is clearly a problem in a real multiprogramming system, where a single CPU core is shared among many processes. Busy waiting also wastes CPU cycles that some other process might be able to use productively.
- The type of mutex lock we have been describing is also called a **Spinlock** because the process “spins” while waiting for the lock to become available.

- **Spinlocks do have an advantage**, however, in that **no context switch is required** when a process must wait on a lock, and a context switch may take considerable time.
- In certain circumstances on multicore systems, spinlocks are in fact the preferable choice for locking. If a lock is to be held for a short duration, one thread can “spin” on one processing core while another thread performs its critical section on another core.
- On modern multicore computing systems, spinlocks are widely used in many operating systems.

# Semaphores

- Semaphore proposed by Dutch computer scientist **Edsger Dijkstra**, is a technique to manage concurrent processes by using a simple integer value, which is known as **Semaphores**.
- A semaphore **S** is an integer variable that, apart from initialization, is accessed only through two standard atomic operations: **wait()** and **signal()**.



## Semaphores (Contd.)

`wait ( )` → **P** [from the Dutch word **proberen**, which means "to test"]

`signal ( )` → **V** [from the Dutch word **verhogen**, which means "to increment"]

- The **wait()** operation was originally termed **P** (from the Dutch *proberen*, “to test”); **signal()** was originally called **V** (from *verhogen*, “to increment”).



## Semaphores (Contd.)

The definition of `wait()` is as follows:

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

The definition of `signal()` is as follows:

```
signal(S) {  
    S++;  
}
```

## Semaphores (Contd.)

- All modifications to the integer value of the semaphore in the **wait()** and **signal()** operations must be executed **atomically**. That is, when one process modifies the semaphore value, no other process can simultaneously modify that same semaphore value.
- In addition, in the case of **wait(S)**, the testing of the integer value of S ( $S \leq 0$ ), as well as its possible modification ( $S--$ ), must be executed without interruption.
- Semaphores are of two types:
  - ❑ **Binary Semaphore:** This is also known as a **mutex lock**, as they are locks that provide mutual exclusion. It can have only two values: **0 and 1**. **Its value is initialized to 1**. It is used to implement the solution of critical section problems with multiple processes and a single resource.
  - ❑ **Counting Semaphore:** Unlike a binary semaphore, which can only take values 0 and 1, a counting semaphore can take any non-negative integer value, which represents the number of available resources. A counting semaphore is initialized with a positive integer value that represents the number of available resources. For example, if a semaphore is initialized to 3, it indicates that there are 3 resources available for use.

- **A critical section is surrounded by both operations to implement process synchronization** (The figure demonstrates the basic mechanism of how semaphores are used to control access to a critical section in a multi-process environment, ensuring that only one process can access the shared resource at a time).

Process P

```
// Some code  
P(s);  
    // critical section  
V(s);  
    // remainder section
```

## Semaphores (Contd.)

- The main disadvantage of the semaphore definition that was discussed is that it requires **busy waiting**.
- While a process is in its critical section, any other process that tries to enter its critical section must loop continuously in the entry code.
- Busy waiting wastes CPU cycles that some other process might be able to use productively.
- This type of semaphore is also called a **spinlock** because the process “spins” while waiting for the lock.

## To overcome the need for busy waiting, we can modify the definition of the wait ( ) and signal ( ) semaphore operations

- When a process executes the wait( ) operation and finds that the semaphore value is not positive, it must wait.

- ☐ However, rather than engaging in busy waiting, the process can block itself.
- ☐ The block operation places a process into a waiting queue associated with the semaphore, and the state of the process is switched to the waiting state.

- Then control is transferred to the CPU scheduler, which selects another process to execute.

# Binary Semaphores

## Down() / Wait () / P ()

```
Down (semaphore s)
{
    if (s.value == 1)
    {
        s.value=0;
    }
    else
    {
        sleep(); // Block this process and place in suspend list
    }
}
```

## Up() / Signal () / V () / Post () / Release ()

```
Up (semaphore s)
{
    if (Suspend List is Empty)
    {
        s.value=1;
    }
    else
    {
        wake up(); // Block to Ready Queue
    }
}
```

# Counting Semaphores

## Down() / Wait () / P ()

```
Down (semaphore s)
{
    s.value = s.value - 1;
    if (s.value < 0)
    {
        sleep(); // Block this process and place in suspend list
    }
    else
        return;
}
```

## Up() / Signal () / V () / Post () / Release ()

```
Up (semaphore s)
{
    s.value = s.value + 1;
    if (s.value <= 0 )
    {
        wake up(); // Block to Ready Queue
    }
}
```

*#Generally FIFO is followed in the Block Queue*

**Question 1:** A counting semaphore S is initialized to 10. Then, 6 P operations and 4 V operations are performed on S. What is the final value of S?



**Question 1:** A counting semaphore S is initialized to 10. Then, 6 P operations and 4 V operations are performed on S. What is the final value of S?

**Solution:**

We know-

P operation also called as wait operation decrements the value of semaphore variable by 1.

V operation also called as signal operation increments the value of semaphore variable by 1.

Thus,

Final value of semaphore variable S

$$= 10 - 6 + 4$$

$$= 8$$

**Question 2:** A counting semaphore S is initialized to 7. Then, 20 P operations and 15 V operations are performed on S. What is the final value of S?

**Question 2:** A counting semaphore S is initialized to 7. Then, 20 P operations and 15 V operations are performed on S. What is the final value of S?

**Solution:**

We know-

P operation also called as wait operation decrements the value of semaphore variable by 1.

V operation also called as signal operation increments the value of semaphore variable by 1.

Thus,

Final value of semaphore variable S

$$= 7 - 20 + 15$$

$$= 2$$

**Question 3:** A shared variable  $x$ , initialized to zero, is operated on by four concurrent processes W, X, Y, Z as follows. Each of the processes W and X reads  $x$  from memory, increments by one, stores it to memory and then terminates. Each of the processes Y and Z reads  $x$  from memory, decrements by two, stores it to memory, and then terminates. Each process before reading  $x$  invokes the P operation (i.e. wait) on a counting semaphore S and invokes the V operation (i.e. signal) on the semaphore S after storing  $x$  to memory. Semaphore S is initialized to two. What is the **maximum** possible value of  $x$  after all processes complete execution?

- A) -2
- B) -1
- C) 2
- D) None of the above

## Solution:

Process W	Process X	Process Y	Process Z
Wait (S)	Wait (S)	Wait (S)	Wait (S)
Read (x)	Read (x)	Read (x)	Read (x)
$x = x + 1;$	$x = x + 1;$	$x = x - 2;$	$x = x - 2;$
Write (x)	Write (x)	Write (x)	Write (x)
Signal (S)	Signal (S)	Signal (S)	Signal (S)

**Solution:**

Process W	Process X	Process Y	Process Z
Wait (S)	Wait (S)	Wait (S)	Wait (S)
Read (x)	Read (x)	Read (x)	Read (x)
$x = x + 1;$	$x = x + 1;$	$x = x - 2;$	$x = x - 2;$
Write (x)	Write (x)	Write (x)	Write (x)
Signal (S)	Signal (S)	Signal (S)	Signal (S)

**Final Answer:  $X=2$**

**Question 4:** A shared variable  $x$ , initialized to zero, is operated on by four concurrent processes W, X, Y, Z as follows. Each of the processes W and X reads  $x$  from memory, increments by one, stores it to memory and then terminates. Each of the processes Y and Z reads  $x$  from memory, decrements by two, stores it to memory, and then terminates. Each process before reading  $x$  invokes the P operation (i.e. wait) on a counting semaphore S and invokes the V operation (i.e. signal) on the semaphore S after storing  $x$  to memory. Semaphore S is initialized to two. What is the **minimum** possible value of  $x$  after all processes complete execution?

- A) -2
- B) -1
- C) 2
- D) None of the above

**Solution:**

Process W	Process X	Process Y	Process Z
Wait (S)	Wait (S)	Wait (S)	Wait (S)
Read (x)	Read (x)	Read (x)	Read (x)
$x = x + 1;$	$x = x + 1;$	$x = x - 2;$	$x = x - 2;$
Write (x)	Write (x)	Write (x)	Write (x)
Signal (S)	Signal (S)	Signal (S)	Signal (S)

**Final Answer:  $X = -4$**



**Question 5:** If a counting semaphore present value is 20, which of the following operations will result in semaphore value 27 ?

A) 3P, 10V, 3V, 2P

B) 8P, 5V, 12V, 2P, 2V

C) 7P, 6V, 5V, 3P, 6V

D) 6P, 2V, 5V, 3P, 6V

**Question 5:** If a counting semaphore present value is 20, which of the following operations will result in semaphore value 27 ?

A) 3P, 10V, 3V, 2P

B) 8P, 5V, 12V, 2P, 2V

C) 7P, 6V, 5V, 3P, 6V

D) 6P, 2V, 5V, 3P, 6V

**Correct Answer: Option C**

**Question 6:** Consider a non-negative counting semaphore  $S$ . The operation  $P(S)$  decrements  $S$ , and  $V(S)$  increments  $S$ . During an execution, 18  $P(S)$  operations and 7  $V(S)$  operations are issued in some order. The largest initial value of  $S$  for which at least one  $P(S)$  operation will remain blocked\_\_\_\_\_?

A) 10

B) 8

C) 9

D) None

**Question 6:** Consider a non-negative counting semaphore S. The operation P(S) decrements S, and V(S) increments S. During an execution, 18 P(S) operations and 7 V(S) operations are issued in some order. The largest initial value of S for which atleast one P(S) operation will remain blocked\_\_\_\_\_?

A) 10

B) 8

C) 9

D) None

$$S - 18 + 7 = -1$$

**Correct Answer: Option A**

# References

1. Abraham Silberschatz, Peter Baer Galvin, Greg Gagne, “Operating System Concepts,” Eleventh Edition (Willey).
2. Andrew S. Tanenbaum, “Modern Operating Systems”, Fourth Edition (Pearson Publications), 2014.
3. <https://www.geeksforgeeks.org/>
4. <https://www.javatpoint.com/>
5. <https://www.tutorialspoint.com/>
6. <https://www.nesoacademy.org/>
7. <https://www.baeldung.com/>
8. <https://www.educative.io/>