

# Operating System

**Unit – 4**

**(Part-A)**

**Deadlock**



**Mayank Mishra**

**School of Electronics Engineering**

**KIIT-Deemed to be University**

# Introduction

- In a multiprogramming environment, several processes may compete for a finite number of resources.
- A process requests resources; if the resources are not available at that time, the process enters a waiting state.
- Sometimes, a waiting process can never again change state, because the resources it has requested are held by other waiting processes. This situation is called a **deadlock**.



*# Perhaps the best illustration of a deadlock can be drawn from a law passed by the Kansas legislature early in the 20th century. It said, in part: “When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up again until the other has gone.”*





## System Model

- ❑ A system consists of a finite number of resources to be distributed among a number of competing processes.
- ❑ The resources may be partitioned into several types (or classes), each consisting of some number of identical instances.
- ❑ CPU cycles, files, and I/O devices (such as network interfaces and DVD drives) are examples of resource types.

## System Model (Contd.)

Under the normal mode of operation, a process may utilize a resource in only the following sequence:

- ❑ **Request:** The process requests the resource. If the request cannot be granted immediately (for example, if a mutex lock is currently held by another process), then the requesting process must wait until it can acquire the resource.
- ❑ **Use:** The process can operate on the resource (for example, if the resource is a printer, the process can print on the printer).
- ❑ **Release:** The process releases the resource.

# Deadlock Characterization

- In a deadlock, processes never finish executing, and system resources are tied up, preventing other jobs from starting.
- A deadlock situation can arise if the following four conditions hold simultaneously in a system:
  - ❑ Mutual exclusion.
  - ❑ Hold and wait.
  - ❑ No pre-emption.
  - ❑ Circular wait.


## Deadlock Characterization (Contd.)

- ***Mutual exclusion:*** At least one resource must be held in a non-sharable mode; that is, only one process at a time can use the resource. If another process requests that resource, the requesting process must be delayed until the resource has been released.
- ***Hold and wait:*** A process must be holding at least one resource and waiting to acquire additional resources that are currently being held by other processes.
- ***No preemption:*** Resources cannot be preempted; that is, a resource can be released only voluntarily by the process holding it, after that process has completed its task.



## Deadlock Characterization (Contd.)

- **Circular wait:** A set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes must exist such that  $P_0$  is waiting for a resource held by  $P_1$ ,  $P_1$  is waiting for a resource held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource held by  $P_n$ , and  $P_n$  is waiting for a resource held by  $P_0$ .
- Let us assume that there are three processes  $P_1$ ,  $P_2$  and  $P_3$ . There are three different resources  $R_1$ ,  $R_2$  and  $R_3$ .  $R_1$  is assigned to  $P_1$ ,  $R_2$  is assigned to  $P_2$  and  $R_3$  is assigned to  $P_3$ .
  - After some time,  $P_1$  demands for  $R_1$  which is being used by  $P_2$ .  $P_1$  halts its execution since it can't complete without  $R_2$ .  $P_2$  also demands for  $R_3$  which is being used by  $P_3$ .  $P_2$  also stops its execution because it can't continue without  $R_3$ .  $P_3$  also demands for  $R_1$  which is being used by  $P_1$  therefore  $P_3$  also stops its execution.
  - In this scenario, a cycle is being formed among the three processes. None of the process is progressing and they are all waiting. The computer becomes unresponsive since all the processes got blocked.





## Resource Allocation Graph (RAG)

- A **Resource Allocation Graph (RAG)** is a visual way to understand how **resources** are assigned in an **operating system**.
- Instead of using only tables to show which resources are allocated, requested, or available, the RAG uses nodes and edges to clearly illustrate relationships between **processes** and their required resources.
- RAGs primarily help in detecting deadlocks by visually representing the relationships between processes and resources, making it easier to identify potential deadlock conditions.

# RAG (Contd.)

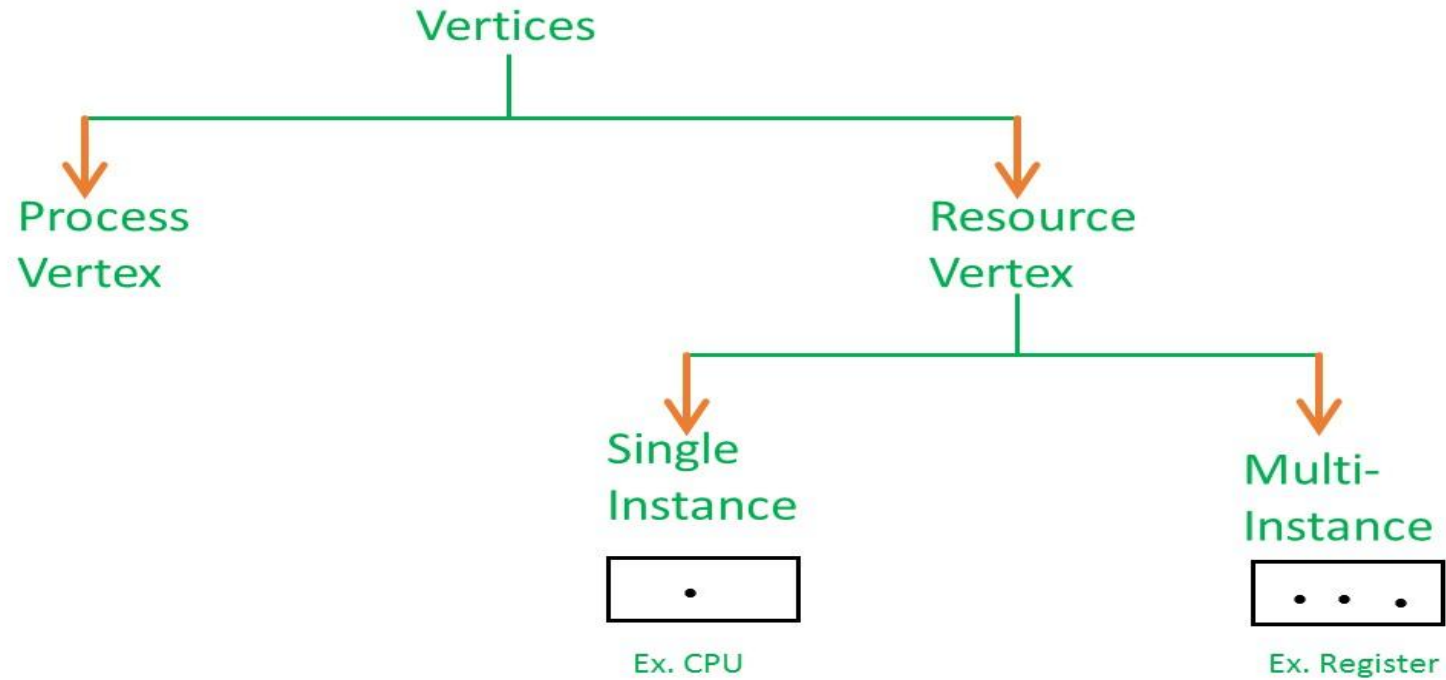
## Types of Vertices in RAG

In a Resource Allocation Graph, there are two types of vertices

- 1. Process Vertex:** Every process will be represented as a process vertex. Generally, the process will be represented with a **circle**.
- 2. Resource Vertex:** Every resource will be represented as a resource vertex. It is also two types:
  - ❑ **Single Instance Type Resource:** A Single Instance Resource refers to a type of resource in the system that has only one available instance or copy. In the Resource Allocation Graph (RAG), a single instance resource is represented as a single resource node, and only one process can be assigned to this resource at any time.
  - ❑ **Multi-Resource Instance Type Resource:** A Multi Instance Resource refers to a type of resource that has multiple instances available. In a Resource Allocation Graph (RAG), a multi-instance resource has multiple resource nodes connected to process nodes, allowing multiple processes to request and hold instances of that resource simultaneously.

# RAG (Contd.)

## Types of Vertices in RAG




# RAG (Contd.)

## Types of Edges in RAG


There are two types of edges in RAG:

- ❑ **Assign Edge:** If you already assign a resource to a process then it is called Assign edge. This is shown by an arrow from the resource vertex to the process vertex.
- ❑ **Request Edge:** A request edge represents that a process is currently requesting a resource. This is shown by an arrow from the process vertex to the resource vertex.



*# If a process is using a resource, an arrow is drawn from the resource node to the process node. If a process is requesting a resource, an arrow is drawn from the process node to the resource node.*


## Example 1 (Single Instances RAG)



**Does it contains Deadlock?**




## Example 1 (Single Instances RAG)



Let's see, whether this RAG has cycle:\_\_\_\_\_?

## Example 1 (Single Instances RAG)



Let's see, whether this RAG has cycle: **YES**


## Example 1 (Single Instances RAG)

- If RAG has circular wait (cycle), it will always have Deadlock.

*# Above statement is only true for Single Instance Scenario*

# Example 1 (Single Instances RA)

*Other approach to check*



Process	Allocation		Request	
	R1	R2	R1	R2
P1	1	0	0	1
P2	0	1	1	0

## Example 1 (Single Instances RAG)

*Other approach to check*

Process	Allocation		Request	
	R1	R2	R1	R2
P1	1	0	0	1
P2	0	1	1	0

Availability = ( R1, R2) = (0,0)

Can the request of P1 and P2 be fulfilled?



## Example 1 (Single Instances RAG)

*Other approach to check*


Process	Allocation		Request	
	R1	R2	R1	R2
P1	1	0	0	1
P2	0	1	1	0

Availability = ( R1, R2) = (0,0)

Can the request of P1 and P2 be fulfilled? - **No**


In this RAG, Deadlock is present.

## Example 2 (Single Instances RAG)



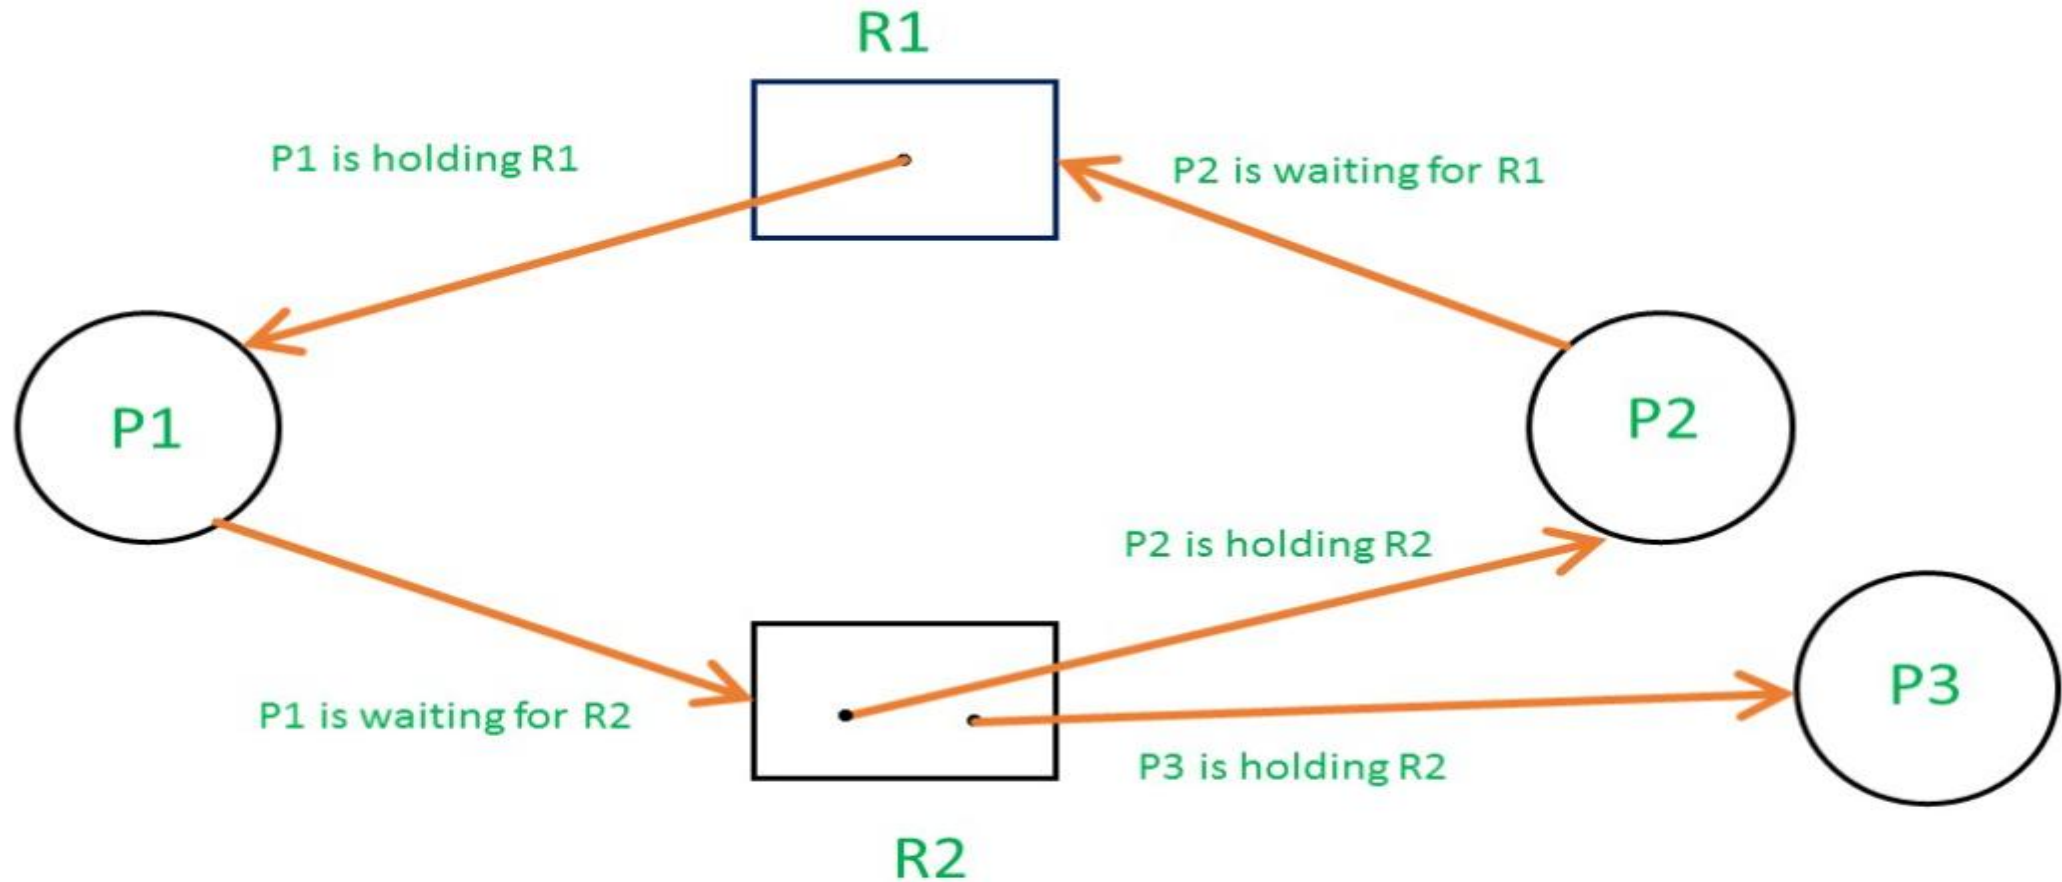
**Does it contains Deadlock?**

## Example 2 (Single Instances RAG)

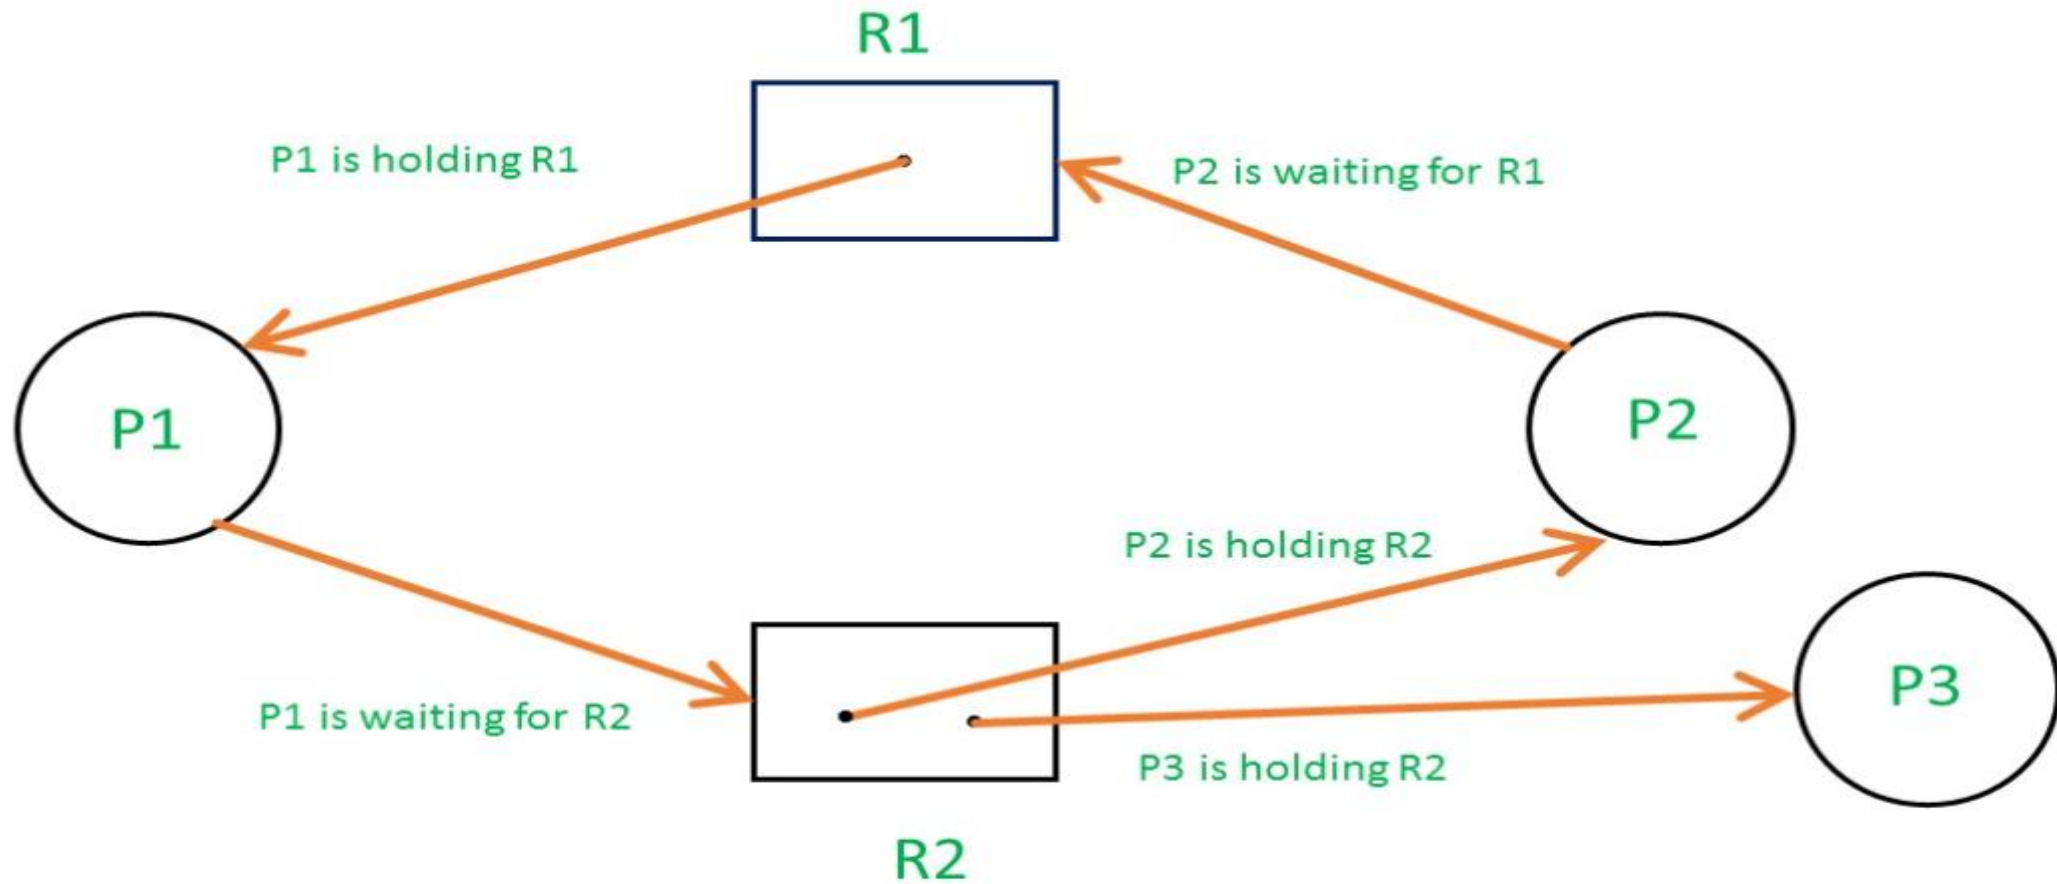


Does it contains Deadlock?-- **NO**

## Example 3 (Multi Instances RAG)

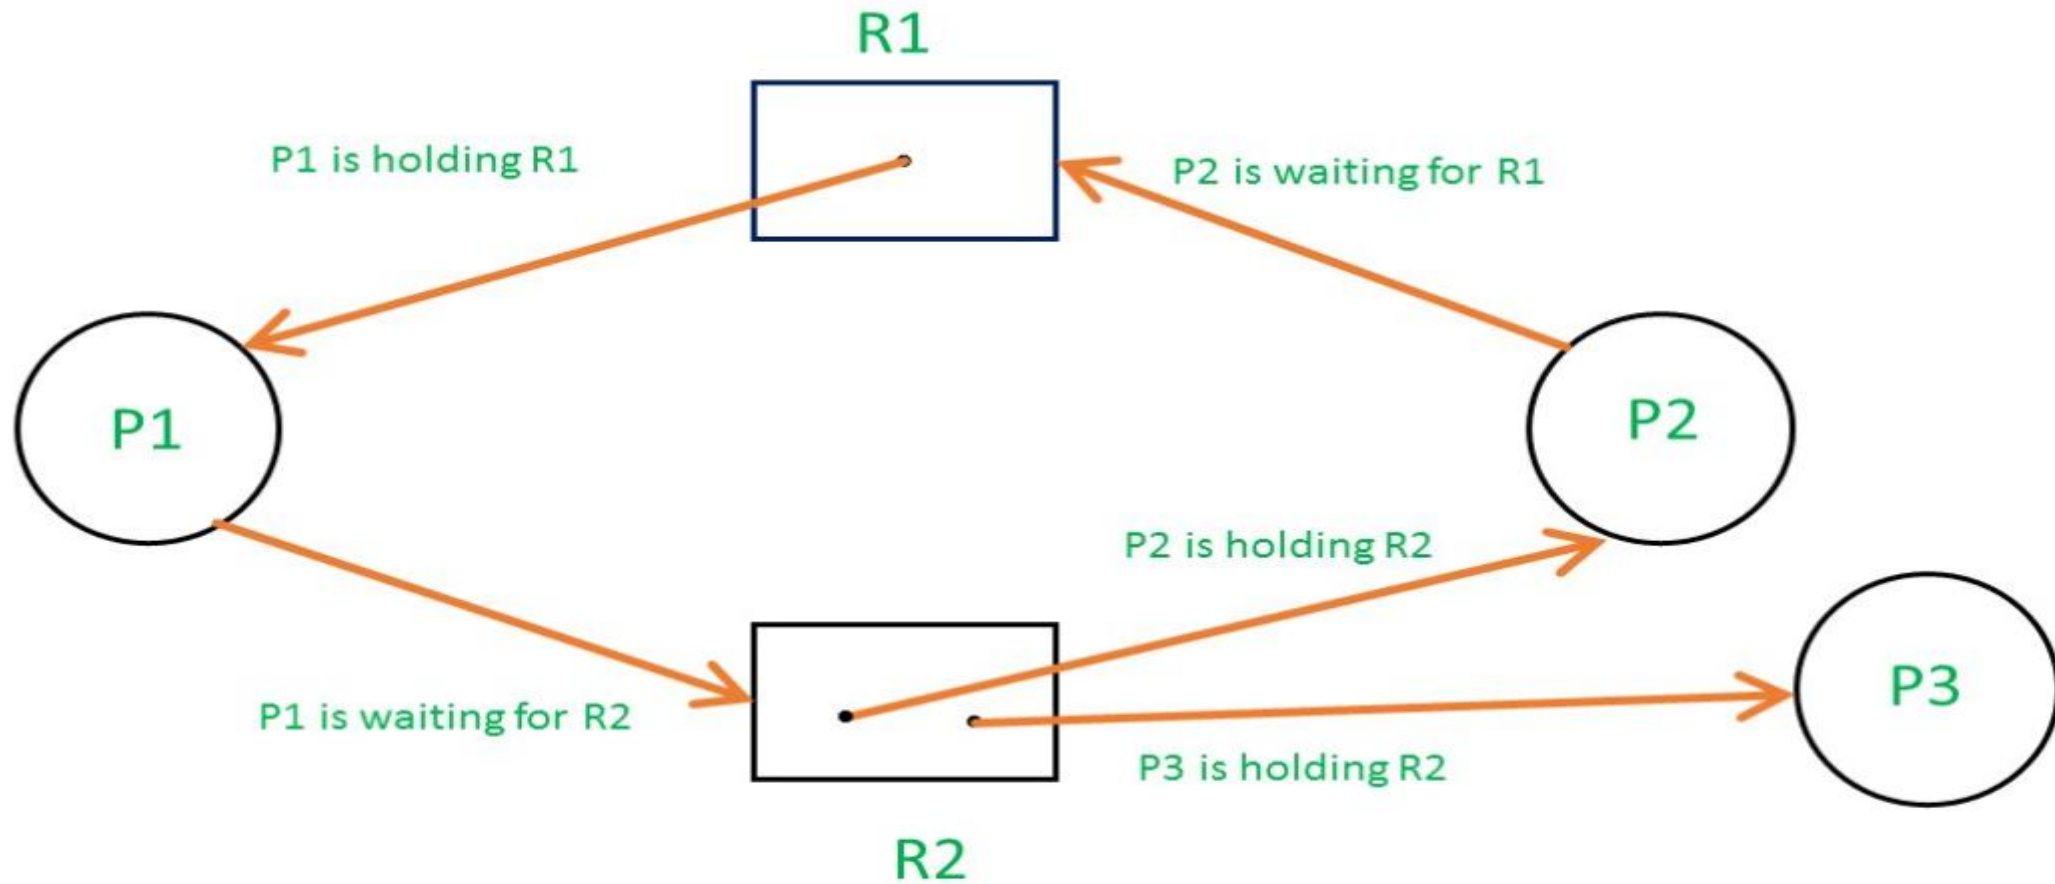


**Does it contains Deadlock?**



**Let's see, whether this RAG has cycle:\_\_\_\_\_?**





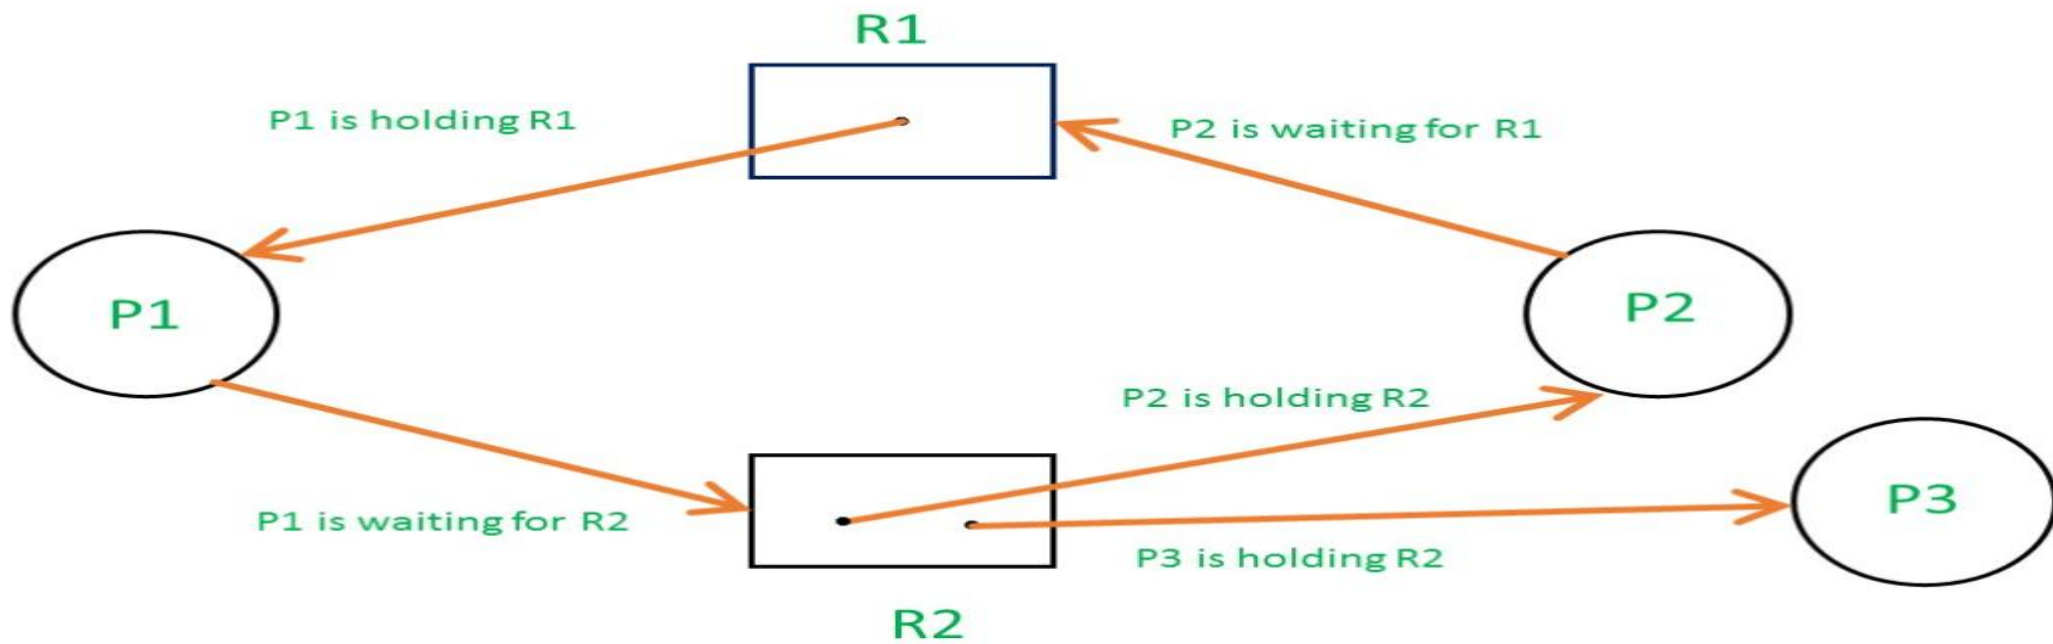
Let's see, whether this RAG has cycle: **YES**

Let's see, whether this RAG has cycle: **YES**

Can there will be Deadlock?

*# As observed in the Single Instance Scenario*

Let's check with the other approach!!!!



Process	Allocation		Request	
	R1	R2	R1	R2
P1	1	0	0	1
P2	0	1	1	0
P3	0	1	0	0


Process	Allocation		Request	
	R1	R2	R1	R2
P1	1	0	0	1
P2	0	1	1	0
P3	0	1	0	0

Available = 0 0 (As P3 does not require any extra resource to complete the execution and after P3 completion P3 release its own resource)  
P3 0 1

New Available = 0 1 (As using new available resource we can satisfy the requirement of process P1  
P1 1 0 and P1 also release its previous resource)

New Available = 1 1 (Now easily we can satisfy the requirement of process P2)  
P2 0 1


New Available = 1 2



No Deadlock


- There is no deadlock in this RAG. Even though there is a cycle, still there is no deadlock. Therefore, in multi-instance resource cycle is not sufficient condition for deadlock.

## Example 4 (Multi Instances RAG)



**Does it contains Deadlock?**






Process	Allocation		Request	
	Resource		Resource	
	R1	R2	R1	R2
P1	1	0	0	1
P2	0	1	1	0
P3	0	1	1	0

Process	Allocation		Request	
	Resource		Resource	
	R1	R2	R1	R2
P1	1	0	0	1
P2	0	1	1	0
P3	0	1	1	0

- So, the Available resource is = (0, 0), but requirement are (0, 1), (1, 0) and (1, 0). So, you can't fulfill any one requirement. Therefore, it is in deadlock. Therefore, every cycle in a multi-instance resource type graph is not a deadlock. If there has to be a deadlock, there has to be a cycle. So, **in case of RAG with multi-instance resource type, the cycle is a necessary condition for deadlock but not sufficient.**

## Example 5 (Multi Instances RAG)



Find if the system is in a deadlock state otherwise find a safe sequence.?

	Allocation			Need		
	R1	R2	R3	R1	R2	R3
Process P0	1	0	1	0	1	1
Process P1	1	1	0	1	0	0
Process P2	0	1	0	0	0	1
Process P3	0	1	0	0	2	0

Available = [ R1 R2 R3 ] = [ 0 0 1 ]

Available

$$= [0\ 0\ 1] + [0\ 1\ 0]$$

$$= [0\ 1\ 1]$$



Available

$$= [0\ 1\ 1] + [1\ 0\ 1]$$

$$= [1\ 1\ 2]$$



Available

$$= [1\ 1\ 2] + [1\ 1\ 0]$$

$$= [2\ 2\ 2]$$



Available

$$= [2\ 2\ 2] + [0\ 1\ 0]$$

$$= [2\ 3\ 2]$$

- Process P2 is allocated the requested resources.
- It completes its execution and then free up the instances of resources held by it.

- Process P0 is allocated the requested resources.
- It completes its execution and then free up the instances of resources held by it.

- Process P1 is allocated the requested resources.
- It completes its execution and then free up the instances of resources held by it.

- Process P3 is allocated the requested resources.
- It completes its execution and then free up the instances of resources held by it.

Thus,

- ❑ There exists a safe sequence P2, P0, P1, P3 in which all the processes can be executed.
- ❑ So, the system is in a safe state.

# References

1. Abraham Silberschatz, Peter Baer Galvin, Greg Gagne, “Operating System Concepts,” Eleventh Edition (Willey).
2. Andrew S. Tanenbaum, “Modern Operating Systems”, Fourth Edition (Pearson Publications), 2014.
3. <https://www.geeksforgeeks.org/>
4. <https://www.javatpoint.com/>
5. <https://www.tutorialspoint.com/>
6. <https://www.nesoacademy.org/>
7. <https://www.baeldung.com/>
8. <https://www.educative.io/>
9. <https://prepinsta.com>